

15.4.4)

$$B(R) = 1000$$

$$B(S) = 500$$

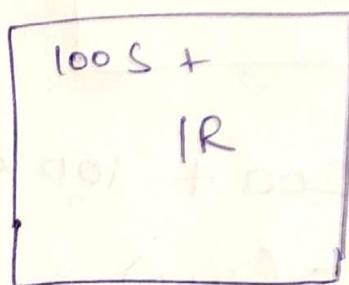
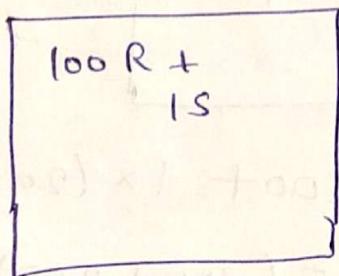
$$M = 101$$

a) Seating cost  $\Rightarrow 4(B(R) + B(S)) = 4(1500)$   
 $= 6000$

Now there are 2 ways of doing nested loop join : To keep 100 blocks of S & 1 of R or 100 of S & 1 of R.

$$M = 101$$

$$M = 101$$



$$2(500 + 5 \times 250)$$

$$= 2(1750)$$

$$= 3500$$

$$2(250 + 3 \times 500)$$

$$= 2(1750)$$

$$= 3500$$

Because we have to bring all of S, 5 times as many of R will have same value.

Same in both

Total Cost  $\Rightarrow 6000 + 3500$   
 $= 9500$

b) Sorting =  $4(1000 + 500) = 6000$

Joining  $\Rightarrow$

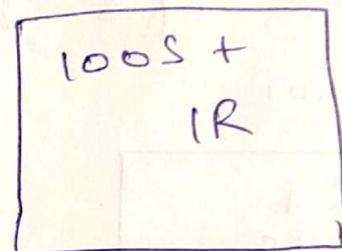
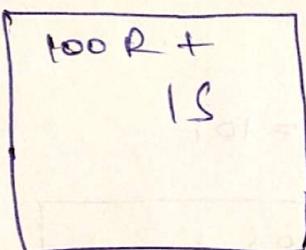
$$R = 1000$$

200	200	200	200	200
-----	-----	-----	-----	-----

$$S = 500$$

100	100	100	100	100
-----	-----	-----	-----	-----

5 yr values each  
equally likely.



$$5(200 + 100 \times 2)$$

$$= 5(400) = 2000$$

$$5(100 + 1 \times (200))$$

$$= 5(100 + 200)$$

$$(200 = 1500)$$

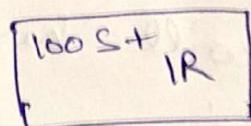
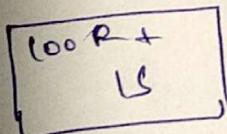
Total Cost  $\Rightarrow 6000 + 1500 = 7500$ .

}

c) Sorting =  $4(1000 + 500) = 6000$

$$R = 100, \dots 10 \text{ times}$$

$$S = 50, \dots 10 \text{ times}$$



$$\Rightarrow 10(100 + 1 \times 50)$$

$$\Rightarrow 10(150) = 1500$$

$$\Rightarrow 5(100 + 200(1)) = 1500$$

or,

$$10(50 + 50 \times 2) = 1500$$

(same)

$\Rightarrow$  Total Cost =  $6000 + 1500$   
 $= 7500$

15.4.5.) In this algorithm we combine second stage sorting with joining.

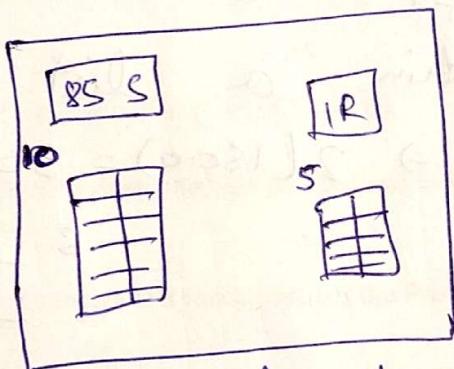
a) So sorting within a list takes 1 read & write  $\Rightarrow 2(1000 + 500) = 3000$

Now half values in each R & S are same.

R =	<table border="1"> <tr> <td>500</td> <td>500</td> </tr> </table>	500	500
500	500		

280	250
-----	-----

General : There will be 10 sorted list of R & 5 of S as 100 is capacity.  
So, we have to make 10 list of R & 5 of S.



By general rule of thumb of nested loop we know that we keep lesser value relation in outer loop. So, S is in outer loop.

Joining

$$\Rightarrow 2 \left( 250 + \left[ \frac{250}{85} \right] \times 500 \right) \Rightarrow 2(250 + 3(500)) \Rightarrow 3500$$

$$\text{So, Total Cost} = 3000 + 3500 \\ = 6500$$

b.) Sorting within a list would take 1 read & write  $\Rightarrow 2(1000 + 500) = 3000$   
 Now 5 values are same.

R.

200	200	200	200	200
-----	-----	-----	-----	-----

100	100	100	100	100
-----	-----	-----	-----	-----

Same 85 S + 1R & others for 1 block from each list.

Joining Cost  $\Rightarrow$

$$5 \left( 100 + \left[ \frac{100}{85} \right] 200 \right)$$

$$\Rightarrow 5 \left( 100 + 2(200) \right) \Rightarrow 2800$$

$$\text{Total Cost} = 3000 + 2800 = 5800$$

c.) Sorting within a list would take 1 read & write  $\Rightarrow 2(1500) = 3000$

R

100	- - - 10 times
-----	----------------

S

50	- - - 10 times
----	----------------

Same as 85 S + 1R & others for 1 block from each list.

Joining Cost  $\Rightarrow$

$$\Rightarrow 10 \left( 50 + \left[ \frac{50}{85} \right] \times 100 \right)$$

$$\Rightarrow 10 \left( 50 + 100 \right) = 1500$$

$$\text{Total Cost} = 3000 + 1500$$

$$= 4500$$

## 15.5.5

a)

Given:

$$B(R) = 1000$$

$$B(S) = 500$$

$$M = 101$$

Also, its given: "to speed up the join, we want to use as few buckets as possible (assuming tuples distribute evenly between buckets)"

Basically, we want to:

1. Use as few buckets as possible
2. Utilize as much of main memory available.
3. During the joining step of hash-join algorithm, fit one complete bucket of a relation into the main memory.

From above statements, we can say that:

$$\frac{500}{k} \leq 101$$

Therefore,  $k \geq 5$  (where k = number of buckets to be used).

Hence, the minimum number of buckets that can be used and satisfy 1,2,3 is k=5.

Now, lets calculate the cost:

R and S are initially stored in consecutive blocks. (as per the Professor, we can assume this)

### Hashing Step:

Total Main Memory Blocks = 101

Reserved blocks: 5 (one for each bucket), 1(for calculating hash values)

Therefore, available blocks =  $101 - (5+1) = 95$  blocks.

### Read R:

We will read R in chunks of 95 blocks.

Hence, the read time is:  $(100 + 95/2) * \text{ceil}(1000/95) = 1622.5 \text{ ms}$

### Write R:

When we write a block to disk, we can specify the address where we want to write it.

Hence, each block is going to be written to disk with a random I/O cost.

Thus, the writing time is:

$$100.5 * 1000 = 100500 \text{ ms}$$

### Read S:

We will read S in chunks of 95 blocks.

Hence, the read time is:  $(100 + 95/2) * \text{ceil}(500/95) = 885 \text{ ms}$

### Write S:

Each block is going to be written to disk with a random I/O cost.

Thus, the writing time is:

$$100.5 * 500 = 50250 \text{ ms}$$

### Joining Step:

During the hashing step, the write were such that all blocks of a bucket got stored consecutively onto disk. So, we can read chunks of blocks of the same bucket.

In this step, we will read into main memory a complete bucket of S, and then read the corresponding bucket of R, one block at a time. We'll repeat this for all buckets.

Hence, the joining time is:

$$\text{Read S: } (100 + 100/2) * 5 = 750 \text{ ms}$$

$$\text{Read R: } (100.5 * 1000) = 100500 \text{ ms}$$

Therefore, total I/O time is:

$$1622.5 + 100500 + 885 + 50250 + 750 + 100500 = 254507.5 \text{ ms} = \text{approx } 254.5 \text{ s}$$

b)

Given:

$$B(R) = 1000$$

$$B(S) = 500$$

$$M = 101$$

Also, its given: "to speed up the join, we want to use as few buckets as possible (assuming tuples distribute evenly between buckets)"

For hybrid hash join, to save some disk I/O's, we keep one COMPLETE bucket of one relation in main memory.

Yet again, the conditions here are similar to the one in the previous part:

We want to:

1. Use as few buckets as possible => make each bucket size as close to the main memory size as possible.
2. Utilize as much of main memory available.
3. During the hashing step of hash-join algorithm, one complete bucket of a relation and one block for rest of the buckets should fit into the main memory.

From the above conditions, we have:

$$\frac{500}{k} \leq 101 \Rightarrow k \geq 5$$

For  $k = 5$ , a bucket size would be  $500/5 = 100$  blocks, and 4 other blocks would be needed for rest of the buckets during hashing step. So the total memory requirement would be  $100 + 4 = 104$  blocks whereas we have only 101 blocks. So, we choose the next higher value of  $k$  i.e.  $k = 6$ .

Now, lets calculate the cost of hybrid hash join:

R and S are stored initially in consecutive blocks respectively (as per the Prof., we can assume this)

### Hashing Step:

Total Main Memory Blocks = 101

Reserved blocks:  $\text{ceil}(500/6) = 84$  (one complete bucket), 5 (one block for rest of the buckets), 1 (for calculating hash)

Therefore, available blocks =  $101 - (84+5+1) = 11$  blocks.

### Read S:

We read S in chunks of 11 blocks.

Hence, the read time is:

$$(100 + 11/2) * \text{ceil}(500/11) = 4853 \text{ ms}$$

### Write S:

We will write back  $500 - \lceil 500/6 \rceil = 500 - 84 = 416$  blocks into disk, as blocks of one complete bucket will keep on residing in main memory.

Furthermore, when we write a block to disk, we can specify the address where to write it. Hence, each block is going to be written to disk with a random I/O cost.

Thus, the writing time is:

$$416 * 100.5 = 41808 \text{ ms}$$

### Read R:

We read R in chunks of 11 blocks:

Hence, the read time is:

$$(100 + 11/2) * \lceil 1000/11 \rceil = 9600.5 \text{ ms}$$

### Write R:

We will write back  $1000 - \lceil 1000/6 \rceil = 1000 - 167 = 833$  blocks into disk, as join output for the corresponding bucket of R, whose S bucket already resides in the main memory, can be done while calculating the hash for blocks of R i.e. we write back 5 of the 6 buckets of R.

Furthermore, when we write a block to disk, we can specify the address where to write it. Hence, each block is going to be written to disk with a random I/O cost.

Thus, the writing time is:

$$833 * 100.5 = 83716.5 \text{ ms}$$

### Joining Step:

During the hashing step, the write were such that all blocks of a bucket got stored consecutively onto disk. So, we can read chunks of blocks of the same bucket.

In this step, we will read into main memory a complete bucket of S, and then read the blocks of the corresponding bucket of R (number of free blocks left after reading a complete bucket of S). We'll repeat this for all buckets.

Total Main Memory Blocks = 101

Blocks for S =  $\text{ceil}(500/6) = 84$  (one complete bucket of S),

Blocks for R =  $101 - 84 = 17$  blocks.

### Read S:

We read S in chunks of 84 blocks (i.e. one complete bucket)

Hence, total read time:

$$(100 + 84/2) * \text{ceil}(416/84) = 710 \text{ ms}$$

### Read R:

We read R in chunks of 17 blocks:

Hence, total read time:

$$(100 + 17/2) * \text{ceil}(833/17) = 5316.5 \text{ ms}$$

Hence, the total I/O time is:

$$4853 + 41808 + 9600.5 + 83716.5 + 710 + 5316.5 = 146004.5 \text{ ms} = \text{approx } 146 \text{ sec}$$

### c) Sort-based join

Given: We write sorted sublists to consecutive blocks of disks.

### Creating sorted sublists:

#### Read R:

We read R in chunks of 100 blocks.

Hence, read time:

$$(100 + 100/2) * \text{ceil}(1000/100) = 1500 \text{ ms}$$

#### Write R:

We write R in chunks of 100 blocks.

Hence, write time:

$$(100 + 100/2) * \text{ceil}(1000/100) = 1500 \text{ ms}$$

#### Read S:

We read S in chunks of 100 blocks.

Hence, read time:

$$(100 + 100/2) * \text{ceil}(500/100) = 750 \text{ ms}$$

### Write S:

We write S in chunks of 100 blocks.

Hence, write time:

$$(100 + 100/2) * \text{ceil}(500/100) = 750 \text{ ms}$$

### Sorting the sorted sublists:

#### Read R:

To maximally utilize the main memory space and the fact that the sublists were written in consecutive blocks, we read 10 blocks of each sorted sublist at a time. Thus, main memory used is =>  $10 * 10 = 100$  blocks

Hence, the read time is:

$$(10 * (100 + 10/2)) * 10 = 10500 \text{ ms}$$

#### Write R:

We write back one block at a time. Hence, the write cost is:

$$1000 * 100.5 = 100500 \text{ ms}$$

#### Read S:

To maximally utilize the main memory space and the fact that the sublists were written in consecutive blocks, we read 20 blocks of each sorted sublist at a time. Thus, main memory used is =>  $20 * 5 = 100$  blocks

Hence, the read time is:

$$5 * (5 * (100 + 20/2)) = 2750 \text{ ms}$$

#### Write S:

We write back one block at a time. Hence, the write cost is:

$$500 * 100.5 = 50250 \text{ ms}$$

## Joining:

During the write step of sorting the sorted sublists, both R and S got stored in consecutive blocks. The minimum joining time is found when, we simply read 100 blocks of R and 1 block of S. This gives us, joining time:

Read R:  $(100 + 100/2) * \text{ceil}(1000/100) = 1500 \text{ ms}$

Read S:  $100.5 * 500 = 50250 \text{ ms}$

Hence, the total I/O time is:

$$1500 + 1500 + 750 + 750 + 10500 + 100500 + 2750 + 50250 + 1500 + 50250 = 220250 \text{ ms} = \text{approx } 220 \text{ s}$$

16.2.8.) a)

" Let's consider a table :

a	b
1	5
1	6
1	5
2	4
2	4

$$\text{L.H.S} \Rightarrow (\forall_{a, \text{sum}(b)} \rightarrow x(R)) \Rightarrow \begin{array}{cc} a & b \\ 1 & 16 \\ 2 & 8 \end{array}$$

$$(\forall_{\min(a) \rightarrow y, n} (\forall_{a, \text{sum}(b)} \rightarrow x(R)) \Rightarrow \begin{array}{cc} a & b \\ 1 & 16 \\ 2 & 8 \end{array})$$

$$R.H.S \Rightarrow (\forall_{\min(a) \rightarrow y, b(R)}) \Rightarrow \begin{array}{cc} a & b \\ 1 & 5 \\ 1 & 6 \\ 2 & 4 \end{array}$$

$$(\forall_{y, \text{sum}(b)} \rightarrow x (\forall_{\min(a) \rightarrow y, b(R)}) \Rightarrow \begin{array}{cc} a & b \\ 1 & 11 \\ 2 & 4 \end{array})$$

So, L.H.S  $\neq$  R.H.S.

They are not equal.

(b) Let's consider a table:

a	b
1	5
2	5
1	6
4	4

$$\text{L.H.S} \Rightarrow (\forall a, \max(b) \rightarrow \mu(R)) \Rightarrow \begin{array}{cc} a & b \\ 1 & 6 \\ 2 & 5 \\ 4 & 4 \end{array}$$

$$(\forall \min(a) \rightarrow y, \mu(\forall a, \max(b) \rightarrow \mu(R))) \Rightarrow \begin{array}{cc} a & b \\ 1 & 6 \\ 2 & 5 \\ 4 & 4 \end{array}$$

$$\text{R.H.S} \Rightarrow (\forall \min(a) \rightarrow y, b(R)) \Rightarrow \begin{array}{cc} a & b \\ 1 & 5 \\ 4 & 4 \end{array}$$

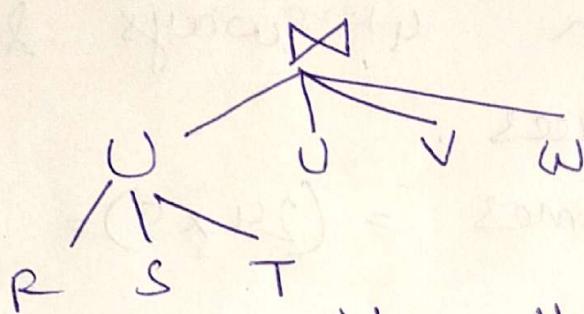
$$(\forall y, \max(b) \rightarrow \mu(\forall \min(a) \rightarrow y, b(R)))$$

$$\Rightarrow \begin{array}{cc} a & b \\ 1 & 6 \\ 4 & 4 \end{array}$$

So, L.H.S  $\neq$  R.H.S.

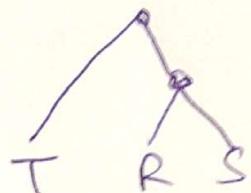
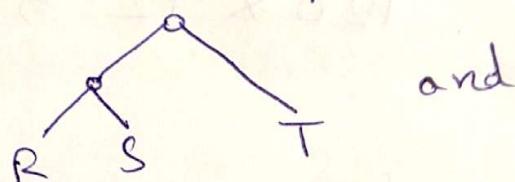
They are not equivalent.

16.3.5.) This is the final tree.



Let us first consider the union part.

Two structures are possible  $\Rightarrow$

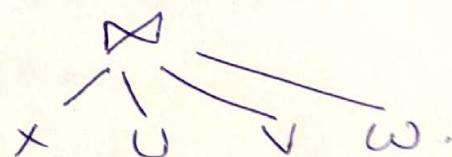


The leaves can be arranged or parenthesizing  
in  $3! = 6$  ways. There are 2 trees.

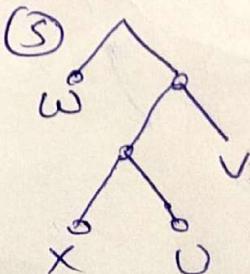
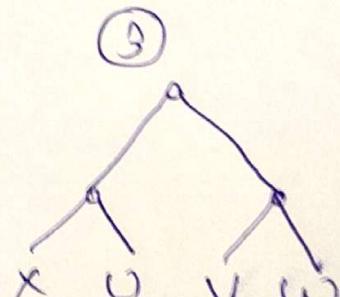
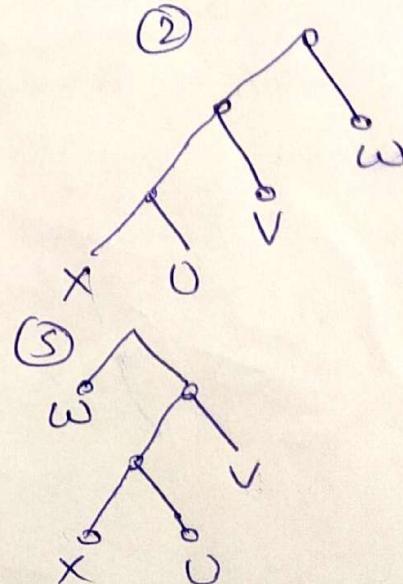
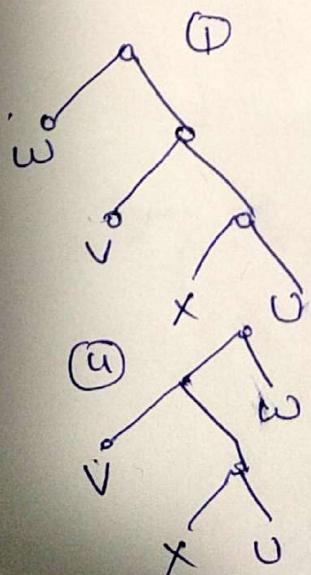
So total ways  $= 6 \times 2 = 12$

Now let us denote   
by X.

So our tree becomes



5 trees are possible  $\Rightarrow$



In each of these leaves can be arranged in  $4!$  ways & there are 5 such trees.

$$\text{So, total becomes } = (24 \times 5) = 120$$

So total exp tree all  $\Rightarrow$

$$120 \times 12 \Rightarrow 1440 \text{ trees}$$

16.4.1.)

$\omega(a, b)$

$T(\omega) = 100$

$V(\omega, a) = 20$

$V(\omega, b) = 60$

$x(b, c)$

$T(x) = 200$

$V(x, b) = 50$

$V(x, c) = 100$

$y(c, d)$

$T(y) = 300$

$V(y, c) = 50$

$V(y, d) = 50$

$z(d, e)$

$T(z) = 400$

$V(z, d) = 40$

$V(z, e) = 100$

(a)  $\omega \bowtie X \bowtie Y \bowtie Z$

$$T(\omega \bowtie X) \Rightarrow \frac{T(\omega) T(X)}{\max\{V(\omega, b), V(X, b)\}} = \frac{100 \times 200}{\max\{60, 50\}}$$

$$\Rightarrow \frac{100 \times 200}{60} \Rightarrow \frac{2000}{6}$$

Now we use preservation of values.

$$\Rightarrow V(\omega, a) = 20 \quad V(X, c) = 100$$

$$V(\omega \bowtie X, b) = \min(60, 50) = 50$$

$$\text{Now, } T(\omega \bowtie X \bowtie Y) = \frac{2000}{6} \times \frac{300}{\max(100, 50)} = \frac{2000 \times 50}{100} = 1000$$

$$T(\omega \bowtie X \bowtie Y \bowtie Z) \Rightarrow \frac{1000 \times 400}{80} = 8000 \Rightarrow \text{Ans.}$$

b)  $T_{a=10}(\omega) \Rightarrow \frac{T(\omega)}{V(\omega, a)}$

$$\Rightarrow \frac{100}{20} = 5 \text{ Ans.}$$

$$c) \quad \sigma_{c=20}(Y) = \frac{T(Y)}{V(Y, c)}$$

$$= \frac{300}{50} = 6 \text{ Ans.}$$

$$d) \quad \sigma_{c=20}(Y) \propto Z$$

$$\text{Now, } \sigma_{c=20}(Y) = \frac{T(Y)}{V(Y, c)} = \frac{300}{50} = 6$$

$$T(\sigma_{c=20}(Y) \times Z) = \frac{T(\sigma_{c=20} Y) \times T(Z)}{\max \left\{ \sum_{d=1}^{\min(6, 40)} V(Y, d), V(Z, d) \right\}}$$

$$\Rightarrow \frac{6 \times 400}{\max(50, 40)} = \frac{6 \times 400}{40} = 60$$

$$e) \quad W \times Y = T(W) \times T(Y) = 100 \times 300 \\ = 30000$$

$$f) \quad \sigma_{d>10}(Z) \Rightarrow \frac{1}{3} T(Z) \quad \text{Assume that } \frac{1}{3} \text{ values are greater than 10.}$$

$$\Rightarrow \left[ \frac{1}{3} (400) \right] = 134$$

$$g) \quad \sigma_{a=1 \text{ AND } b=2}(W) = \frac{T(W)}{V(W, a) \times V(W, b)}$$

$$= \frac{100}{20 \times 60} = \frac{1}{12} = 1 \text{ tuple}$$

We consider  $\frac{1}{12}$  as 1 as otherwise none of conditions could be satisfied.

h)

$$a=1 \text{ AND } b > 2 (\omega)$$

$\Rightarrow$  We assume  $b$  has value greater than 2.

$$\Rightarrow \frac{T(\omega)}{\sqrt{(\omega, a)} \times 3} \Rightarrow \frac{10^5}{10^5 \times 3} = \left\lceil \frac{5}{3} \right\rceil = 2 \text{ tuples.}$$

i)

$$X \bowtie X.c < Y.c Y$$

In this case do cartesian product followed by condition. So we assume only third satisfies that condition.

$$\Rightarrow \frac{T(X) \times T(Y)}{3} \Rightarrow$$

$$= \frac{200 \times \frac{10^5}{200}}{3} = 20000$$

7) 16.5.7.) On this question we have to prove that if E  $\bowtie$  F are not on top, then we would get an optimal value. So let us consider an example of a query plan.

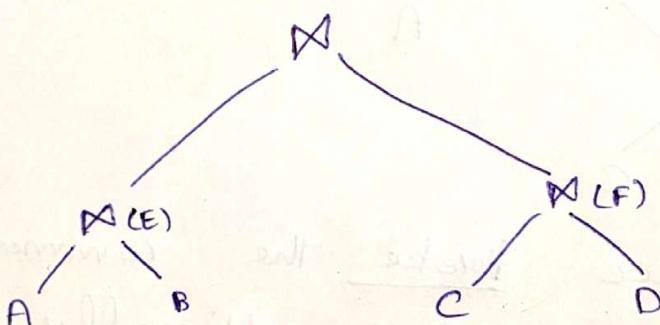
$A \bowtie B \bowtie C \bowtie D$

$$E = A \bowtie B$$

$$F = C \bowtie D$$

Optimal plan of E is only  $A \bowtie B$

Optimal plan for F is only  $C \bowtie D$



Now we have to make an example where we first combine E & F.

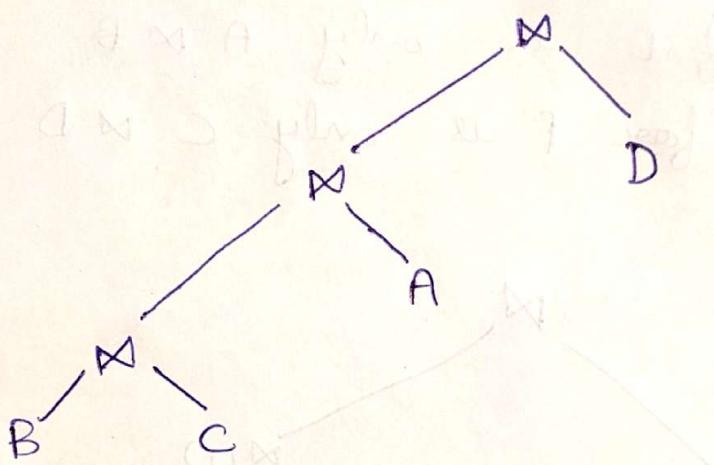
Let us consider that B & C do not have any common value for join attribute. There can be many examples for that.

- 1) B could be  $\sigma_{d=100} X$   
C could be  $\sigma_{d=50} Y$   
where d is join attribute.

- 2) B could be  $\sigma_p > 500 X$   
C could be  $\sigma_p < 300 X$   
p is join attribute.

Now if we join  $B \bowtie C$  first, then  
we are left with deletion with 0 tuples.  
Now combining relation with 0 tuple with  
 $A \bowtie D$  would result in minimum cost.

Exp tree :



So, this way we believe the common idea  
that calculating  $E \bowtie F$  optimally & then  
joining  $E \bowtie F$  at top would result in  
optimal query plan.